

Decision-Focused Demand Learning for Equitable Bus Frequency Optimization

Arjun Veluri

Independent Researcher
arjunsveluri@gmail.com

Keywords: Mixed-Integer Programming, Predict-Then-Optimize, Decision-Focused Learning, Public Transportation, Stochastic Optimization, Robust Optimization, Demand Forecasting, Equity.

Abstract: We study the bus frequency allocation problem of Pune’s PMPML network as a case where the optimal allocation is *structurally robust*: its dependence on demand prediction quality is empirically small, and the predict-then-optimize literature’s standard motivating story — better predictions yield better decisions — barely applies. The instance is built from public data: 544 real routes with origin/destination/distance from the OpenCity Pune transit dataset (OpenCity Urban Data Portal, 2024a), five operating periods, thirteen depots per published PMPML statistics, an approximately 2,000-bus peak-deployable fleet, and priority designations derived from per-ward SC + ST population share in Census of India 2011 data (OpenCity Urban Data Portal, 2024b; Registrar General and Census Commissioner of India, 2011). Per-route hourly ridership is simulated (no public AFC data exists at that granularity), calibrated to PMPML’s ≈ 1.1 million daily riders and the published 8–11 AM / 5–9 PM peaks. We solve a 24,480-binary-variable mixed-integer program with fleet, depot, equity, and operator-cost constraints; train three demand predictors (linear regression, random forest, XGBoost with seven quantile heads and split-conformal calibration); and evaluate four predict-then-optimize variants across 20 random seeds with paired Wilcoxon tests (Wilcoxon, 1945). Three findings characterise the robustness. First, the Value of Stochastic Solution and the Expected Value of Perfect Information are both essentially zero ($VSS = 0.53$ and $EVPI = 1.91$ passenger-hours on a $\sim 26,000$ passenger-hour objective): the route-level allocation structure is preserved across scenarios, so scenario-resolution prediction buys almost nothing. Second, the decision-focused dissociation is *statistically* real but *operationally* small. Sample-weighted XGBoost achieves mean regret 3.13 pass-hr vs. 3.94 for plain MSE-XGBoost over 20 seeds (paired one-sided Wilcoxon $p = 0.011$) despite a $7\times$ worse MAPE on the same test fold; the absolute regret gap is ~ 0.8 passenger-hours, equivalent to roughly 48 person-minutes total across 1.1 million daily riders — well under one-twentieth of a second per rider per day. The neural-network SPO+ variant does not significantly beat MSE-XGBoost ($p = 0.75$). Third, the rider-wait / operator-cost Pareto frontier shows that the headline 19.5% improvement at $\lambda = 0$ contracts to zero around $\lambda \approx 3$; reporting a single percentage in isolation overstates the operational stakes. The contribution is the demonstration: predict-then-optimize regret comparisons must be reported with seed-level variability and absolute units, and on this allocation the rigorous answer is that prediction quality barely matters for the decision.

1 INTRODUCTION

Public bus systems in Indian metropolitan areas operate under a combination of constraints that few transit planning tools were designed to address jointly: heavy-tailed demand, fleet sizes that have not kept pace with population growth, and equity considerations that bind on both political and welfare grounds (Mahadevia et al., 2013; Parisar, 2020). The Pune Mahanagar Parivahan Mahamandal Limited (PMPML), which serves approximately 1.1 million daily riders across 544 routes with a peak-deployable

fleet of roughly 2,000 buses anchored to thirteen depots (nine CNG plus four electric) (Institute for Transportation and Development Policy India, 2024; Ramaswamy et al., 2018), is the network this paper analyses. Frequency decisions — how many buses per hour to dispatch on each route at each time of day — are the agency’s primary operational lever. Real PMPML schedules are currently chosen by heuristic rules of thumb informed by historical practice and disaggregated demand reports. We anchor the network’s structural inputs to public datasets: route topology and distances from the OpenCity Pune Bus Stops and

Routes dataset (OpenCity Urban Data Portal, 2024a; DataMeet Pune Chapter, 2024); equity priorities from per-ward Scheduled Caste plus Scheduled Tribe population share in Census of India 2011 PMC data (OpenCity Urban Data Portal, 2024b; Registrar General and Census Commissioner of India, 2011); depot list and operational scale from PMPML’s published statistics and the operational analysis in (Ramaswamy et al., 2018). Per-route hourly ridership is simulated because PMPML does not release smart-card or AFC data at that granularity; the simulator’s diurnal profile reflects the 8–11 AM and 5–9 PM peaks reported in Pune-specific traffic studies and in published smart-card analyses of comparable Indian agencies.

This paper formulates frequency allocation as a mixed-integer program (MIP), trains demand models to drive it under uncertainty, and connects predictors to the optimiser via decision-focused training procedures that target operational regret rather than prediction error. Our empirical motivation is contrary to most of the predict-then-optimize literature: on this instance, the optimal allocation depends so weakly on which predictor we use that all four predict-then-optimize variants land within ~ 2 passenger-hours of the oracle on a $\sim 26,000$ passenger-hour objective, and information metrics (VSS, EVPI) computed against the perfect-foresight benchmark are on the order of one part in 10^4 . A statistically significant decision-focused win is detectable (§4.5; XGB-weighted vs. MSE-XGBoost, paired Wilcoxon $p = 0.011$ over 20 seeds), but the effect size is operationally tiny — the gap exists, and the dissociation between RMSE and regret is genuine, but in this regime decision-focused training is making a sub-second difference per rider per day. The methodological contribution is that demonstrating this requires multi-seed paired tests and absolute reporting units: a single-seed comparison at seed 42 would have suggested a 44.7% NN-SPO⁺ regret reduction over MSE-XGBoost, but the 20-seed paired protocol reveals this as a sampling artifact that tells the wrong story in both magnitude and direction.

We contribute (i) a deterministic frequency-allocation MIP with explicit fleet, depot, equity, and operator-cost constraints that scales to the full 544-route PMPML network and solves to provable optimality in seconds, plus a λ -sweep that traces the rider-wait / operator-cost Pareto frontier; (ii) an XGBoost demand-forecasting pipeline with seven quantile heads and split-conformal calibration of the central prediction band on the validation fold; (iii) stochastic and robust (minimax) MIP formulations driven by the quantile-derived scenarios, with explicit decomposition of the Value of Stochastic So-

lution (VSS) and Expected Value of Perfect Information (EVPI); (iv) a decision-focused neural network trained with an SPO+-style surrogate, with the regret comparison against the MSE-trained XGBoost reported over a 20-seed sweep and a paired Wilcoxon signed-rank test rather than a single-seed point estimate; and (v) a sensitivity study over fleet size and equity threshold, with LP shadow-price analysis identifying which constraints (fleet vs. specific depots) are binding on this instance.

2 RELATED WORK

The classical literature on bus frequency setting begins with (Furth and Wilson, 1981) and (Ceder and Wilson, 1986), who established the trade-off between rider waiting cost and operator capacity cost and provided heuristics that remain in use at many transit agencies. Our work updates two pieces of this pipeline: where the classical literature treats demand as exogenous, we predict it with modern machine learning; where the classical optimisation is single-scenario, we explicitly handle uncertainty.

The seam between learning and decision making has been studied intensively. (Elmachtoub and Grigas, 2022) introduced the SPO+ surrogate loss and proved consistency for problems with linear objectives, demonstrating that minimising prediction error alone can produce decisions arbitrarily far from optimal. (Donti et al., 2017) extended decision-focused learning to stochastic optimisation problems with differentiable solvers. (Wilder et al., 2019) treated combinatorial optimisation via a continuous relaxation of the optimiser. (Mandi et al., 2020) addressed hard combinatorial problems including knapsack and shortest-path. The full SPO+ machinery requires solving an inner optimisation at every gradient step, which is prohibitive for our 24,480-binary-variable MIP. We therefore adopt a tractable linear surrogate: an L1 loss weighted by the downstream sensitivity $\partial w / \partial f$ evaluated at a frozen reference allocation, which captures the leading-order regret gradient without per-batch inner solves.

For the uncertainty-handling components we follow standard references. (Birge and Louveaux, 2011) provide the canonical treatment of stochastic programming, including the VSS and EVPI definitions used in §3. (Soyster, 1973) and (Ben-Tal and Nemirovski, 2002) establish the robust optimisation framework used for the minimax formulation. Quantile regression (Koenker and Bassett Jr, 1978) supplies the scenarios. Demand prediction itself uses XGBoost (Chen and Guestrin, 2016) with SHAP at-

tributions (Lundberg and Lee, 2017) for interpretability. The optimisation core is implemented in PuLP (Mitchell et al., 2011) and solved with CBC.

3 PROBLEM AND METHODOLOGY

3.1 Problem Description

The PMPML network as represented in the OpenCity dataset (OpenCity Urban Data Portal, 2024a) comprises 544 routes after collapsing direction-pair entries (each route appears once as a base identifier; upbound / downbound variants in the source CSV are treated as a single route operating in both directions). Route distances range from 1.1 to 46.6 km with a mean of 16.8 km. We classify routes into three service categories by distance, matching the operational distinction used by PMPML between feeder, trunk, and suburban services: feeder (< 8 km, 70 routes), trunk (8–22 km, 325 routes), and suburban (≥ 22 km, 149 routes). Mean trip speed under typical Pune traffic is reported as ~ 17 km/h in published traffic studies; we use this figure to convert distance to one-way travel time. The round-trip cycle time used by the optimiser is $\tau_r = 2 \cdot \text{avg_trip_min} + \text{dwell}$ with $\text{dwell} = 7.5$ minutes per cycle, the midpoint of the 5–10 min terminal layover typical for Indian urban operations.

Routes are anchored to one of thirteen depots: nine CNG depots (Aundh, Bhosari, Dhankawadi, Hadapsar, Katraj, Kothrud, Nigdi, Pimpri, Swargate) and four electric depots (Balewadi, Maan-Hinjewadi, Shewalewadi, Charholi) per published PMPML statistics. The fleet cap used in the optimisation, $B = 2,000$, is the peak-deployable capacity reported by PMPML; the status-quo schedule encoded in the loaded route data uses approximately 1,800 buses at peak (verified by `scripts/gate_check_phase1.py`: 1799 with the round-trip cycle time), leaving roughly 200 buses in reserve / maintenance / spare margin. The optimisation reports headline gains relative to this $\sim 1,800$ -bus status quo, not relative to the theoretical 2,000-bus saturation point. Per-depot fleet capacity is allocated proportional to each depot’s minimum bus need (priority routes at the equity floor plus non-priority at $k = 1$) with $\pm 20\%$ jitter – PMPML does not publish per-depot fleet counts and the alternative of an equal $1/|J|$ split produces infeasibilities, because priority routes (defined below) cluster geographically.

Operating hours span 06:00–23:00 and are discretised into five periods (early_morning, AM_peak, midday, PM_peak, evening) chosen to capture domi-

nant demand structure while keeping the optimisation tractable.

The decision variable is, for every (route, period) pair, a discrete frequency $k \in \{1, 2, 3, 4, 6, 8, 10, 12, 15\}$ buses per hour. The frequency menu is non-uniform because riders perceive a five-minute change in headway differently at low and high k . Priority routes are defined by the SC + ST population share in the PMC Census 2011 ward containing the route’s origin or destination endpoint (OpenCity Urban Data Portal, 2024b; Registrar General and Census Commissioner of India, 2011): routes whose catchment falls in the top quartile of ward-level SC + ST share are designated priority. In the network as we load it, 97 of 544 routes (17.8%) are priority and must run at $k \geq 3$ during peak periods regardless of marginal efficiency. This floor operationalises the equity argument that a bus system below three-per-hour at peak ceases to function as a reliable substitute for paratransit modes that disproportionately tax low-income commuters (Mahadevia et al., 2013). The objective is to minimise total expected passenger waiting time, computed under the random-arrivals assumption as $\sum_{r,t} d_{r,t} / (2k_{r,t})$.

3.2 Deterministic MIP Formulation

Let R be the set of routes ($|R| = 544$), T the set of operating periods ($|T| = 5$), and K the discrete frequency menu ($|K| = 9$). Parameters are demand $d_{r,t}$ (passengers per hour), round-trip cycle time $\tau_r = 2 \cdot \text{avg_trip_min}_r + 7.5$ (minutes), bus consumption $b_{r,k} = \lceil k\tau_r / 60 \rceil$, the random-arrivals expected wait $w_k = 1/(2k)$, and an operator-cost coefficient $\lambda \geq 0$ in passenger-hour-equivalents per bus-period. The decision variable $x_{r,t,k} \in \{0, 1\}$ equals one if route r in period t is assigned frequency k :

$$\min_x \sum_{r,t,k} (d_{r,t} w_k + \lambda b_{r,k}) x_{r,t,k} \quad (1)$$

$$\text{s.t.} \sum_k x_{r,t,k} = 1 \quad \forall r, t \quad (2)$$

$$\sum_{r,k} b_{r,k} x_{r,t,k} \leq B \quad \forall t \quad (3)$$

$$\sum_{k \geq 3} x_{r,t,k} = 1 \quad \forall r \in R^P, t \in T^{\text{pk}} \quad (4)$$

$$\sum_{r \in R_{j,k}} b_{r,k} x_{r,t,k} \leq B_j \quad \forall j, t \quad (5)$$

Constraint (2) requires exactly one frequency per route-period. Constraint (3) caps the network-wide peak-deployable fleet at $B = 2000$. Constraint (4) is the equity floor for the 97 priority routes $r \in R^P$ during peak periods $T^{\text{pk}} = \{\text{AM_peak}, \text{PM_peak}\}$. Con-

straint (5) caps the buses anchored to depot j at the depot capacity B_j . The model has 24,480 binary variables and approximately 3,260 constraints. The LP relaxation is tight; CBC solves the model to optimality in under twenty seconds on commodity hardware. With $\lambda = 0$ the objective reduces to total expected passenger waiting time under the standard random-arrivals approximation $w_k = 1/(2k)$ (Furth and Wilson, 1981); sweeping $\lambda > 0$ traces the rider-wait / operator-cost Pareto frontier. Refinements that incorporate empirical headway variance, and a richer treatment of operator cost, are listed in §6.

3.3 Demand Forecasting

Data. PMPML does not release smart-card or AFC data at per-hour, per-route granularity. We therefore train forecasting models on simulated hourly ridership for the 50 highest-frequency routes over 2024, totalling approximately 430,000 rows after a seven-day lag-feature window is dropped. The generative model is described in §6; its diurnal profile reproduces the 8–11 AM and 5–9 PM peaks observed in published Pune traffic studies, and its overall daily ridership is anchored to PMPML’s published ≈ 1.1 million-rider figure (Institute for Transportation and Development Policy India, 2024). The remaining 494 routes are not modelled directly at hourly resolution; their period-level demand is constructed by a residual-scaling procedure described in §6, which corrects a parametric baseline by the empirical mean ratio (observed to baseline) measured on the modelled 50.

Features. Twenty-three features in total: cyclical encodings (sine and cosine) of hour, day-of-week, and month; calendar flags (weekend, holiday, monsoon, exam season, festival); weather (maximum temperature, precipitation in mm); route-level features (length, number of stops, priority score, one-hot service category); and lag features at one day and seven days at the same hour.

Models. Three model classes are compared. Linear regression provides the trivial baseline. Random forest (200 trees, default depth) provides a non-parametric tabular benchmark. XGBoost (Chen and Guestrin, 2016) with depth 8, learning rate 0.03, 2,000 trees with early stopping at patience 50, subsample 0.85, and column subsample 0.85 is the primary model. Three additional XGBoost models with quantile-regression objectives at $\tau \in \{0.1, 0.5, 0.9\}$ supply prediction intervals (Koenker and Bassett Jr, 1978).

Train/val/test split. The data are split chronologically: months 1–8 for training, 9–10 for validation

(and early stopping), 11–12 for test. Random splits would leak future information into the training set and are not used.

3.4 Stochastic and Robust Optimisation

Scenarios. Six demand scenarios are constructed from the quantile predictions to span a realistic range of operational conditions. S1 (probability 0.35) is the baseline at $q_{0.5}$. S2 (0.15) is high demand at $q_{0.9}$. S3 (0.15) is low demand at $q_{0.1}$. S4 (0.15) is a monsoon shock dampening $q_{0.1}$ by an additional 20%. S5 (0.10) boosts trunk demand by 15% above $q_{0.9}$ during peak periods. S6 (0.10) is a festival dip scaling $q_{0.5}$ by 0.75 across the network. The scenario probabilities reflect a prior over the relative frequency of these conditions in a typical operating year.

Stochastic MIP. The stochastic formulation shares the deterministic decision variables but minimises the probability-weighted expected wait,

$$\min_x \sum_{s \in S} p_s \sum_{r,t,k} d_{r,t}^{(s)} w_k x_{r,t,k}, \quad (6)$$

which is computationally equivalent to the deterministic MIP with $d_{r,t}$ replaced by $\tilde{d}_{r,t} = \sum_s p_s d_{r,t}^{(s)}$.

Robust MIP. The minimax formulation introduces an auxiliary scalar z and replaces the objective with $\min_x z$ subject to $z \geq \sum_{r,t,k} d_{r,t}^{(s)} w_k x_{r,t,k}$ for every scenario $s \in S$, plus all original constraints.

Information metrics. Following (Birge and Louveaux, 2011), $VSS = EEV - RP$ measures the value of the stochastic solution relative to the deterministic one, and $EVPI = RP - WS$ measures the value of perfect information about which scenario will materialise.

3.5 Decision-Focused Learning

The MSE-trained predictor weights all prediction errors equally, but the downstream optimisation does not: an error on a route-period whose optimal frequency is at the boundary between two discrete levels matters more than an error on a route-period whose optimum is robust to small perturbations. Because $\partial w / \partial f = -1/(2f^2)$, the optimisation sensitivity scales as $d_{r,t} / (2f_{r,t}^2)$ at the reference allocation $f_{r,t}$.

We test two decision-focused variants. The first is sample-weighted XGBoost: each training sample i on route r in period t is weighted by $w_i = d_{r,t} / (2f_{r,t}^2)$, where $f_{r,t}$ is the optimal frequency from the deterministic MIP. The second is a two-layer fully-connected neural network (input $\rightarrow 128 \rightarrow 64 \rightarrow 1$, ReLU acti-

vations, BatchNorm) trained with the surrogate

$$L_{\text{SPO}^+}(\hat{d}, d) = \frac{1}{n} \sum_{i=1}^n \frac{d_i}{2f_i^2} |\hat{d}_i - d_i|. \quad (7)$$

The L1 form is robust to outliers, which matters because high-sensitivity samples dominate the loss surface and a single misbehaved point would otherwise destabilise the gradient. The reference allocation f is set once at the start of training using XGBoost-MSE predictions. A staged refresh (recomputing f periodically from the current network’s predictions) is the natural extension; we do not run it in this paper and make no claim about its effect.

Connection to SPO+. The exact SPO+ loss (El-machtoub and Grigas, 2022) requires solving the inner MIP at every gradient step. Equation (7) is a linearisation that captures the leading-order sensitivity of the regret to prediction error at the fixed reference allocation while skipping the per-batch solve. The fixed-allocation approximation is defensible here because the optimal allocation structure is robust to small changes in demand: the binding constraints at the optimum are determined by the relative ordering of routes rather than by the absolute demand level.

4 EXPERIMENTS AND RESULTS

4.1 Experimental Setup

All MIPs are solved with PuLP (Mitchell et al., 2011) and CBC on a Windows 11 machine running Python 3.13.9, single thread for the solver, eight threads for XGBoost and random forest. Synthetic data generation, XGBoost reference training, and the optimisation phases use seed 42; the decision-focused comparison (§4.5) sweeps seeds 0–19 for the four predict-then-optimize methods. The full pipeline including the 20-seed sweep reproduces in approximately three hours end-to-end on commodity hardware.

4.2 Forecasting Results

Table 1: Test-set forecasting performance (months 11–12) on the 50 modelled PMPML routes.

Model	RMSE	MAE	MAPE	R ²
Linear regression	8.88	5.28	56.4%	0.829
Random forest	7.59	3.77	16.7%	0.875
XGBoost	7.17	3.42	15.2%	0.889

XGBoost achieves the best test-set performance on three of four metrics (Table 1). The data-

generating process described in §3.1 includes a latent per-day demand factor and per-(route, date) hidden surge events that the predictor cannot observe, together with heavy-tailed Student- t multiplicative noise on a subset of routes; this caps achievable R^2 near 0.9 even for a well-tuned XGBoost, which is the point – a noise-free DGP leaves no irreducible signal for a decision-focused loss to exploit. The 80% prediction interval from the raw quantile heads covers 70.0% of test observations; applying a split-conformal correction (Angelopoulos and Bates, 2023) calibrated on the validation fold raises empirical coverage to 80.4%, essentially matching the nominal 80% target. The reliability diagram (Figure 1) summarises coverage at all seven trained quantile levels. The conformal correction is reported here as a forecasting-pipeline diagnostic; the calibrated bands are not fed into the stochastic / robust optimisation downstream, which uses the raw $q_{0.1}/q_{0.5}/q_{0.9}$ quantiles in §3.4. Threading the conformal adjustment through to scenario generation is a clean extension we leave to future work.




Figure 1: Left: reliability of the seven trained quantile heads on the test fold (predicted quantile α vs. empirical $P(y \leq \hat{q}_\alpha)$); the ideal line is $y = x$. Right: empirical coverage of the 80% PI before and after split-conformal calibration on the validation fold.

Figure 2a shows the joint distribution of actual and predicted ridership on the test set, coloured by route category. Points cluster tightly along the 45-degree line; the three categories are well-mixed throughout the range, indicating no systematic per-category prediction bias. Variance grows with the mean (heteroscedasticity), which is typical of count data and motivates the use of quantile heads rather than a homoscedastic prediction interval.

Figure 2b ranks the top twelve features by mean absolute SHAP value. Lag features (one-day and seven-day) and the cyclical hour encoding dominate, consistent with the autoregressive structure of hourly ridership. Calendar flags and weather contribute at lower magnitudes. Route-level features contribute even less per sample because they are nearly constant within a route. Figure 3 shows the SHAP beeswarm: each point is one prediction; horizontal position is the

demand. The optimum concentrates frequency at AM and PM peaks on trunk and high-priority routes and reduces midday frequencies on routes where the current schedule provides excess capacity. Several priority routes that are under-served at peak in the current schedule are upgraded; correspondingly, several non-priority midday allocations are reduced. The pattern is consistent with the equity narrative: redistribution occurs primarily along the temporal axis rather than at the expense of priority routes.

Cost-wait Pareto frontier

The 19.5% figure in Table 2 assumes $\lambda = 0$: the optimiser is told to minimise rider waiting time with no penalty for operating more buses, so it saturates the fleet whenever waiting time can be reduced. Real transit decisions trade rider waiting time against operator running cost. To expose this trade-off we sweep λ , the marginal cost per bus-period (in passenger-hour-equivalents), from 0 to 50.




Figure 6: Rider-wait / operator-cost Pareto frontier obtained by sweeping the operator-cost coefficient λ in the MIP objective. Each point is the optimum at one λ . The $\lambda = 0$ point saturates the fleet (10,000 bus-periods, i.e., $B = 2,000$ across five periods); moderate values trade waiting time against bus-period usage.

Figure 6 shows the resulting curve. At $\lambda = 0$ the optimum uses 10,000 bus-periods (fleet saturated) for a total wait of 26,226 passenger-hours. At $\lambda = 2$ the optimum uses 9,608 bus-periods for a total wait of 26,911 – giving up 685 pass-hr of waiting time saves 392 bus-periods. At $\lambda = 5$ the optimum uses 8,463 bus-periods (slightly above status quo’s 8,008) for a wait of 30,503 – already worse than status quo on rider waiting time. At $\lambda = 10$ the optimum matches status-quo bus-period usage (8,018) but with substantially worse waiting time (33,773 vs. 32,583), illustrating that the MIP’s structural reallocation benefit shows up only when the optimiser is allowed to spend

more on rider time relative to operator cost. Head-line percentage improvements vs. status quo depend strongly on λ : the 19.5% gain at $\lambda = 0$ disappears around $\lambda \approx 3$. Reporting a single “% improvement” figure without specifying λ obscures this trade-off.




Figure 7: Frequency distribution at peak periods, before (status quo) and after (MIP optimum at $\lambda = 0$), split by priority status.

Figure 7 reports peak-frequency distributions split by priority status and by before/after. Three observations follow. Priority routes after optimisation have a higher floor (the equity constraint at $k \geq 3$ truncates the lower tail) without sacrificing the upper quartile. Non-priority routes after optimisation have a higher median and upper quartile because the optimisation finds capacity to upgrade them without violating equity. Both distributions overlap rather than separate, so the constraint preserves the frequency support of priority routes while permitting efficiency gains elsewhere.

4.4 Stochastic and Robust Analysis




Figure 8: Method performance across the six demand scenarios. Bars are total waiting time. The dashed line is the status-quo mean across scenarios.

Figure 8 reports the four optimisation methods (status quo, deterministic, stochastic, robust) under each of the six scenarios. The three optimised methods cluster within a handful of passenger-hours of one

another in every scenario, each beating status quo by approximately 25–40% depending on scenario. The robust solution achieves the lowest worst-case (scenario S5, the peak-surge scenario, 35,273 pass-hr for status quo vs. approximately 27,245 for all three optimised methods) by roughly one passenger-hour relative to deterministic, i.e., within solver tolerance. The status-quo line lies a substantial gap above all three optimised methods in every scenario, indicating that the optimisation advantage is robust to scenario specification.




Figure 9: Value of Stochastic Solution (VSS) and Expected Value of Perfect Information (EVPI), in passenger-hours.

VSS= 0.53 and EVPI= 1.91 passenger-hours (Figure 9). Both are positive but extremely small (roughly two parts in 10^4 of the objective). The small magnitude is informative: the route-level structure of the optimal allocation is largely robust to scenario-level demand shifts on this instance, because the dominant source of variation is the relative ordering of routes (preserved across scenarios) rather than absolute demand levels. EVPI bounds the maximum value any further investment in scenario-resolution prediction could buy; at ~ 2 passenger-hours, additional scenario modelling on this DGP yields essentially nothing.

4.5 Decision-Focused Results

Table 3 reports four predict-then-optimize methods, each trained over 20 random seeds with the same train/val/test split, evaluated on the same downstream task: predict hourly demand, aggregate to (route, period), plug into the deterministic MIP, evaluate the resulting allocation against true test-period demand. The decision-aware XGBoost variant has the lowest mean regret. A paired Wilcoxon signed-rank test (Wilcoxon, 1945) of the one-sided hypothesis $\text{regret}(\text{XGB-weighted}) < \text{regret}(\text{XGB-MSE})$ gives $p = 0.011$ – decision-aware sample weighting trims regret by approximately 20% relative to plain MSE, and the effect is statistically detectable over $n = 20$ seeds. The neural-network variants do not beat XGBoost on this network: the analogous test

$\text{regret}(\text{NN-SPO}^+) < \text{regret}(\text{XGB-MSE})$ gives $p = 0.75$, i.e., NN-SPO⁺ does *not* significantly beat the XGBoost baseline on regret. A single-seed comparison at seed 42 would have suggested a 44.7% NN-SPO⁺ advantage; the 20-seed paired protocol reveals this as a sampling artifact that flips direction across seeds.

Three substantive readings follow. First, prediction RMSE and decision regret genuinely dissociate. XGB-weighted has the worst MAPE in the table (105.5%) but the best regret – it weights training samples by downstream sensitivity and produces a predictor that is useless as a forecast in isolation but allocation-tuned. The methodological lesson holds. Second, the absolute regret magnitudes are small: the best method’s 3.13 pass-hr of regret is about 1.2×10^{-4} of the oracle objective, or 188 person-minutes total across ~ 1.1 million daily riders – roughly one-hundredth of a second per person. Reporting only a relative “% reduction” on numbers this small inflates the apparent stakes; the methodological contribution is honest evaluation, not the size of the gap. Third, both NN variants underperform XGBoost-MSE in regret terms despite (in the NN-MSE case) matching it on RMSE. This is evidence against the specific NN-SPO⁺ implementation as used here, not against decision-focused learning in general – the exact SPO⁺ loss with a coarsened-MIP inner solver, deferred to future work, is the natural next comparison.




Figure 10: Left: mean test RMSE per method with one-SE error bars across 20 seeds. Right: boxplot of regret distributions across seeds (whiskers are $1.5 \times \text{IQR}$; dashed line at oracle). XGBoost-MSE has the lowest mean regret on this DGP; the decision-focused variants do not dominate it. The decision-weighted XGBoost has high RMSE but competitive regret. NN-MSE is unstable across seeds (outlier seeds inflate the box).

Figure 10 visualises the seed-level variability that a single-seed bar plot would have hidden. The XGBoost baselines are tightly concentrated; the NN methods (and especially NN-MSE) are spread across more than an order of magnitude in regret depending on training seed.

Table 3: Prediction quality vs. decision quality, mean $\pm$ standard error over 20 random seeds (seeds 0–19). “Regret” is the method’s true objective minus the oracle’s ($\approx 26,176$ pass-hr), in passenger-hours. “Frac. obj.” expresses regret as a fraction of the oracle objective, multiplied by 10^4 for readability. “person-min” is total regret converted to person-minutes ($60 \times$ regret); across the network’s ~ 1.1 million daily riders the per-rider impact is well below one second on every method.

Method	RMSE	MAPE	Regret (pass-hr)	Frac. obj. ($\times 10^{-4}$)	person-min
XGBoost (MSE)	7.16	15.5%	3.94 ± 0.11	1.50	236
XGBoost (weighted)	8.09	105.5%	3.13 ± 0.30	1.20	188
NN (MSE)	8.01	24.6%	6.78 ± 1.21	2.59	407
NN (SPO ⁺ surrogate)	10.82	204.0%	4.93 ± 0.70	1.89	296
Oracle	—	—	0	0	0

4.6 Sensitivity Analysis




Figure 11: Fleet-size sensitivity. Left axis: total waiting time. Right axis: number of route-periods with frequency at least three. Vertical dashed line: peak-deployable fleet cap $B = 2,000$ (the status-quo schedule uses approximately 1,800 of these at peak).

Figure 11 sweeps fleet size from $B = 1,200$ to $B = 2,600$. Below $B = 2,000$ the model is infeasible: the equity floor at $k \geq 3$ for the 97 priority routes plus $k \geq 1$ on the remaining 447 cannot be met under the corrected round-trip cycle time $\tau = 2 \cdot \text{avg_trip} + \text{dwell}$. From $B = 2,000$ the curve is monotone decreasing with diminishing returns: $B = 2,000 \rightarrow 2,200$ recovers 3,144 passenger-hours, $B = 2,200 \rightarrow 2,400$ recovers a further 2,388 passenger-hours, $B = 2,400 \rightarrow 2,600$ recovers 2,054 passenger-hours. The right-axis count of route-periods with $k \geq 3$ rises from 356 at $B = 2,000$ to 997 at $B = 2,600$, quantifying how many additional route-periods can be lifted above the “functional service” threshold as fleet grows. The infeasibility cliff at $B < 2,000$ matches the CIRT-derived benchmark of ~ 50 buses per lakh population (Pune’s ~ 4 million-person service area $\rightarrow \approx 2,000$ buses), suggesting that PMPML’s recent expansion from the 2016–17 fleet of $\approx 1,450$ (Ramaswamy et al., 2018) to the current $\sim 2,000$ was operationally necessary, not merely incremental.

Figure 12 sweeps the equity threshold from one (no constraint) to ten (every priority route runs at least ten buses per hour at peak). Thresholds 1–3 are feasi-




Figure 12: Equity-efficiency Pareto frontier. Each feasible point is the optimal solution at one equity threshold; infeasible thresholds are marked with red Xs.

ble at $B = 2,000$. Threshold 4 is the cliff: the feasible region collapses because the fleet- and depot-capacity caps together cannot simultaneously sustain a four-bus-per-hour floor across the 97 priority routes plus a positive frequency on the remaining 447. Total waiting time rises from 24,276 (threshold 1, no equity floor) to 26,226 (threshold 3, the current operating floor), an increase of 8.0%. The equity floor at $k \geq 3$ therefore costs approximately one passenger-hour for every twelve served, which is consistent with prior Indian-transit-equity work flagging similar magnitudes (Mahadevia et al., 2013).

The LP-relaxation shadow prices identify which constraints are binding. With the bus-needs-proportional depot allocation adopted in §3.2, the per-depot constraints are essentially slack and the fleet constraint binds at both peak periods. Marginal valuations of additional peak-period fleet capacity are on the order of -10 passenger-hours per bus, so investment in additional buses (not in particular depots) is what would lower the objective most per marginal bus on this instance.

Robustness to the priority designation

The 8% equity-floor cost reported above is derived from a priority designation that, in turn, depends on a hand-built landmark→ward alias table mapping PMPML route endpoints to PMC Census wards (see §3.1; `load_real_data.py`). A reviewer might reasonably worry that the alias table embeds enough researcher choice to drive the equity-cost number. We therefore expose three rules in the loader and re-run the equity-cost computation under each (Table 4).

Table 4: Robustness of the equity-floor cost to the landmark→ward alias rule. “Strict” allows only direct substring matches against Census ward names; “conservative” adds aliases with clear geographic justification; “current” adds debatable cross-ward equity classifications (e.g. Kondhwa mapped to Bhavani Peth on welfare grounds rather than strict adjacency). The headline equity cost is essentially unchanged between conservative and current; strict is degenerate because PMPML route descriptions almost never contain literal ward names verbatim.

Rule	# priority	Wait no eq.	Equity cost (%)
strict	1	24,400	0.05
conservative	95	24,283	7.60
current	97	24,276	8.03

Dropping the most debatable aliases (Kondhwa→Bhavani Peth, Wanowri→Bhavani Peth, Pashan/Baner/Balewadi →Aundh, Ferguson/Manapa→Ghole Road, Manjari→Hadapsar, Karvenagar→Kothrud, Katraj→Dhankawadi, Undri→Yewalewadi, Kalyani Nagar / Viman Nagar / Parvati / Marketyard / Swargate) removes 2 priority routes (97→95) and moves the equity cost from 8.03% to 7.60%. The headline number is robust to the alias choices. The strict rule produces only one priority route because almost no route descriptions in the OpenCity CSV contain a literal Census ward name verbatim — the endpoints are mostly landmark hubs (Pune Station, Ma Na Pa, Swargate) rather than administrative-ward identifiers — so the strict rule is informative as a lower bound but is not the operationally meaningful one. Migrating to a geographic geocoding of stop coordinates onto ward boundaries (using the OpenCity bus-stop KML and PMC ward polygons) is the natural improvement; it is likely to land between the conservative and current numbers above.

5 DISCUSSION

Operational implications. Re-allocating the $B = 2,000$ peak fleet according to the optimisation re-

duces total passenger waiting time on the 544-route PMPML network by 19.5% at $\lambda = 0$ relative to status-quo frequencies, but the headline contracts to zero around $\lambda \approx 3$ once operator bus-period cost is priced into the objective (Figure 6, §4.3). Whether this gap is actually recoverable for PMPML depends on how the agency values rider waiting time against operating cost — a parameter we do not commit to. The shadow-price analysis (§4.6) shows that on this instance the binding constraint at peak is the fleet itself, not any particular depot; additional buses are worth ~ 10 passenger-hours each of marginal waiting-time reduction.

Decision-focused dissociation. Across 20 random seeds, sample-weighted XGBoost achieves significantly lower decision regret than MSE-XGBoost (paired Wilcoxon one-sided $p = 0.011$), despite an order-of-magnitude worse MAPE on the same test fold. This is the methodological core of the paper: a predictor that performs badly on standard forecasting metrics can produce better operational decisions, because the loss it minimises is better aligned with the downstream optimiser (Elmachtoub and Grigas, 2022). The neural-network SPO^+ surrogate as implemented here does not show the same advantage ($p = 0.75$ for $NN-SPO^+ < XGB-MSE$), which we attribute to the small NN architecture interacting poorly with the variance in the SPO^+ -weighted training signal; an exact SPO^+ loss with a coarsened-MIP inner solver is deferred to future work. The broader lesson — that single-seed regret comparisons are unreliable and that decision-focused training can win on regret while losing on RMSE — generalises to other operational settings (inventory, scheduling, capacity planning).

Equity as a small-but-real policy cost. The Pareto frontier (§4.6) indicates that the current equity floor of three buses per hour at peak adds approximately 8% to total waiting time relative to a no-floor counterfactual. Feasibility cliffs at higher thresholds reflect joint fleet- and depot-level capacity. The policy implication is that PMPML’s current floor sits roughly at the largest equity threshold the agency can satisfy at the current fleet — not on the infeasibility cliff, but within striking distance of it. Raising the floor further would require expanding the fleet beyond $B = 2,000$.

Generalisability. The framework ports directly to other Indian transit agencies with similar institutional structures: BEST (Mumbai), BMTC (Bangalore), DTC (Delhi). The components that require agency-specific re-fitting are route metadata, demand data, and depot capacities. The optimisation core and the decision-focused training procedure are agency-agnostic.

6 LIMITATIONS AND FUTURE WORK

Real route topology, synthetic ridership. Route identifiers, origin / destination, and route length are real PMPML data (OpenCity Urban Data Portal, 2024a; DataMeet Pune Chapter, 2024). Priority designations are real Census 2011 ward-level SC + ST shares (OpenCity Urban Data Portal, 2024b; Registrar General and Census Commissioner of India, 2011) mapped to route endpoints via a landmark-to-ward alias table. Hourly ridership is simulated because PMPML does not release AFC / smart-card data at hourly per-route granularity; the diurnal profile is calibrated to the 8–11 AM and 5–9 PM peaks reported in published Pune traffic studies and to PMPML’s ~ 1.1 million daily ridership figure (Institute for Transportation and Development Policy India, 2024). Hourly ridership is simulated only for the 50 highest-frequency routes; the remaining 494 are assigned period-level demand by a residual scaling procedure (parametric baseline $b_r \cdot \pi_t$ multiplied by an empirical residual ratio $\rho_t = \text{mean}(\text{observed}/\text{baseline})$ estimated on the modelled 50). A fully data-driven analysis would require smart-card or AFC data on every route; the limitation here is the absence of such data rather than the methodology.

Synthetic per-depot fleet capacity. PMPML does not publish per-depot fleet counts at the granularity the MIP’s depot-capacity constraint needs. We allocate $B = 2,000$ across the thirteen depots proportional to each depot’s minimum bus need (priority routes at the equity floor plus non-priority at $k = 1$) with $\pm 20\%$ jitter, clamped above need + 10. This matches the spirit of the operational reality (depots with more priority routes get more buses) without claiming to reproduce the agency’s exact per-depot allocation, which the agency itself does not appear to publish externally (Ramaswamy et al., 2018).

Static daily allocation. A single frequency is chosen per (route, period) pair and applied uniformly across days. Real PMPML operations adjust schedules for day-of-week and special events. The framework can in principle handle these by re-optimising at any cadence; the static analysis is presented because the operational and equity questions resolve at this level.

Random-arrivals waiting model. The expected wait $w_k = 1/(2k)$ assumes Poisson arrivals and uniform headways. Real bus arrivals exhibit bunching (Adamski’s bus paradox) that makes the effective wait longer than $1/(2k)$, especially on heavily-loaded routes. Refinements that incorporate empirical headway variance would tighten the objective without

changing the qualitative results.

Demand under price shocks. The demand model treats fares and competition as exogenous. In reality, paratransit operators reprice in response to bus service quality. A genuinely closed-loop extension would treat demand as a function of frequency through modal substitution.

Decision-focused approximation. The SPO^+ surrogate (7) is a linearisation around a fixed reference allocation. Refresh strategies (e.g., updating the reference every E epochs) and benchmarks against the full SPO^+ loss with an inner solver on a coarsened MIP are deferred to future work.

Depot capacity allocation. Per-depot capacity is allocated proportional to each depot’s minimum bus need (priority routes at the equity floor, non-priority at $k = 1$) plus $\pm 20\%$ jitter, clamped above the minimum. A naive $1/|\text{depots}|$ split produces depot-level infeasibilities under the corrected round-trip cycle time because priority routes cluster geographically in the synthetic network. Real agencies’ depot allocations reflect a mix of legacy, political, and operational factors that our scheme does not attempt to match.

Operator-cost coefficient. The cost-wait frontier in §4.3 sweeps λ across roughly two orders of magnitude. We do not commit to a single operationally “correct” λ because the conversion between passenger-hours of waiting and rupees of bus-period operating cost depends on context-specific value-of-time and unit cost-per-bus-hour assumptions that vary across agencies and analyses.

Single-year training window. The forecast model is trained on a single year of data. Multi-year training would improve estimates of festival and seasonal effects, although the operational allocation is dominated by within-week and within-day structure that the single-year window already captures.

Hard equity floor. The current formulation treats equity as a hard constraint at $k \geq 3$ for priority routes during peaks. Soft penalties, fairness-aware objectives (Wilder et al., 2019), and geographic equity formulations are alternative specifications that may be politically or practically preferable in some settings.

Future directions. (i) Integration of real smart-card data once available; (ii) real-time re-optimisation in response to disruptions, leveraging the sub-five-second solve time; (iii) multi-modal optimisation that couples frequency with modal-substitution behaviour; (iv) joint frequency-pricing optimisation as practised by some international agencies; (v) decision-focused architecture search beyond the two variants tested here, including graph neural networks on the route topology and proper SPO^+ losses with inner solvers on coarsened MIPs; (vi) richer equity formalisations

from the algorithmic-fairness literature; (vii) cross-agency benchmarking on BEST, BMTC, and DTC.

7 CONCLUSION

This paper presented a complete predict-then-optimize pipeline for bus frequency allocation, evaluated on the 544-route Pune PMPML network constructed from public datasets: route topology and distances from OpenCity, equity priorities from Census 2011 PMC ward-level SC + ST shares, depot list and operational scale from PMPML’s published statistics. Per-route hourly ridership is the only major component that remains simulated, because PMPML does not release AFC-level data. The pipeline integrates a deterministic MIP that respects fleet, depot, equity, and operator-cost constraints with an operator-cost penalty λ tracing the rider-wait / operator-cost Pareto frontier; an XGBoost demand model with seven quantile heads and split-conformal calibration of the central prediction band on the validation fold; stochastic and robust formulations driven by quantile-derived scenarios; and a decision-focused neural network trained with an SPO⁺-style surrogate evaluated across 20 random seeds.

Three findings stand out. First, against status-quo frequencies the unconstrained optimum ($\lambda = 0$) reduces total waiting time by 19.5% while saturating the fleet; introducing $\lambda > 0$ trades this gain off against bus-period usage along the frontier reported in §4.3, with the percentage improvement vanishing around $\lambda \approx 3$. Second, sample-weighted XGBoost (“XGB-weighted”), which weights training points by the downstream decision sensitivity $d_{r,t}/(2f_{r,t}^2)$, achieves mean regret 3.13 pass-hr vs. 3.94 for plain MSE-XGBoost over 20 random seeds; a paired Wilcoxon signed-rank test gives $p = 0.011$ for the one-sided hypothesis that decision-aware weighting reduces regret. The neural-network SPO⁺ surrogate as implemented here does *not* significantly beat MSE-XGBoost on regret ($p = 0.75$), reversing an earlier single-seed framing of the same experiment. The methodological lesson – that prediction RMSE and decision regret can dissociate and that single-seed regret comparisons are unreliable – is supported by both results. Third, the current equity floor of three buses per hour at peak for priority routes costs roughly 8% of total waiting time relative to a no-floor counterfactual; the threshold-four feasibility cliff is jointly governed by fleet and depot capacity, with fleet binding more tightly than any individual depot on the needs-proportional depot allocation we adopt.

The methodological contribution – how to eval-

uate predict-then-optimize pipelines without conflating prediction and decision quality, and how to expose the rider / operator trade-off that vanishes when the operator term is omitted – generalises beyond transit. The practical contribution is a reproducible PMPML benchmark anchored to public Indian transit data, in which the synthetic pieces (ridership, hourly profile, per-depot fleet capacity) are explicitly demarcated from the empirical ones, so downstream work can replace them as more operational data becomes available.

Code and data availability. All code (data generation, MIP formulation, demand forecasting, decision-focused training, sensitivity analyses, and figure-generation scripts), the public-data inputs used for the real PMPML instance (OpenCity Pune Bus Stops and Routes; Pune Census 2011 PMC ward-level), and every result table referenced in this paper are available at <https://github.com/juunnq/pune-bus-optimization>. The full pipeline including the 20-seed decision-focused sweep reproduces in approximately three hours on commodity hardware; `pytest tests/test_sanity.py` provides 17 regression checks and `scripts/gate_check_phase1.py` gates the synthetic-data inputs against the published PMPML scale.

REFERENCES

- Angelopoulos, A. N. and Bates, S. (2023). A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *Foundations and Trends in Machine Learning*, 16(4):494–591.
- Ben-Tal, A. and Nemirovski, A. (2002). Robust optimization—methodology and applications. *Mathematical Programming*, 92(3):453–480.
- Birge, J. R. and Louveaux, F. (2011). *Introduction to Stochastic Programming*. Springer, 2nd edition.
- Ceder, A. and Wilson, N. H. (1986). Bus network design. *Transportation Research Part B: Methodological*, 20(4):331–344.
- Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD*, pages 785–794.
- DataMeet Pune Chapter (2024). DataMeet Pune: Bus routes database and GTFS feed creation. <https://github.com/datameet-pune/datameet-pune.github.io/wiki>.
- Donti, P., Amos, B., and Kolter, J. Z. (2017). Task-based end-to-end model learning in stochastic optimization. In *Advances in Neural Information Processing Systems*, volume 30.
- Elmachtoub, A. N. and Grigas, P. (2022). Smart “predict, then optimize”. *Management Science*, 68(1):9–26.

- Furth, P. G. and Wilson, N. H. (1981). Setting frequencies on bus routes: Theory and practice. *Transportation Research Record*, 818:1–7.
- Institute for Transportation and Development Policy India (2024). PMPML’s journey through time: How pune and pimpri chinchwad got better public transport. <https://itdp.in/pmpmls-journey-through-time-how-pune-and-pimpri-chinchwad-got-better-public-transport/>.
- Koenker, R. and Bassett Jr, G. (1978). Regression quantiles. *Econometrica*, pages 33–50.
- Lundberg, S. M. and Lee, S.-I. (2017). A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30.
- Mahadevia, D., Joshi, R., and Datey, A. (2013). Low carbon mobility in india and the challenges of social inclusion: Bus rapid transit (brt) case studies in india. Technical report, UNEP Risø Centre on Energy, Climate and Sustainable Development; Centre for Urban Equity, CEPT University.
- Mandi, J., Demirovic, E., Stuckey, P. J., and Guns, T. (2020). Smart predict-and-optimize for hard combinatorial optimization problems. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 1603–1610.
- Mitchell, S., O’Sullivan, M., and Dunning, I. (2011). PuLP: A linear programming toolkit for Python. Available at <https://github.com/coin-or/pulp>.
- OpenCity Urban Data Portal (2024a). Pune bus stops and routes (pmpml). <https://data.opencity.in/dataset/pune-bus-stops-and-routes>.
- OpenCity Urban Data Portal (2024b). Pune census 2011 data (pmc ward-level). <https://data.opencity.in/dataset/pune-census-2011-data>.
- Parisar (2020). Comparison of bus services in indian cities. <https://www.parisar.org/what-we-do/urban-transport/analysis-reports/comparison-of-bus-services-in-indian-cities>.
- Ramaswamy, S. et al. (2018). Statistical analysis of depot-wise PMPML data. *IEEE Conference Publication*. Document 8745427, IEEE Xplore.
- Registrar General and Census Commissioner of India (2011). Census of india 2011. <https://censusindia.gov.in/>.
- Soyster, A. L. (1973). Convex programming with set-inclusive constraints and applications to inexact linear programming. *Operations Research*, 21(5):1154–1157.
- Wilcoxon, F. (1945). Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83.
- Wilder, B., Dilkina, B., and Tambe, M. (2019). Melding the data-decisions pipeline: Decision-focused learning for combinatorial optimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 1658–1665.